

Forecasting residual analysis

EC 361–001

Prof. Santetti

Spring 2024

Materials

Required readings:

- Hyndman & Athanasopoulos, ch. 6
 - sections 6.3–6.4.

Motivation

Motivation

After introducing the four **benchmark** forecasting models, we will focus on what is **left over** after fitting a forecasting model.

Analogous to a **regression model**, even the best models **cannot** explain the entire variation exhibited by a variable of interest.

In our case, the **fitted/estimated** values of a variable (\hat{y}_t) reflect our **best effort** to predict the future values of a variable of interest.

More formally, a forecasting model's **residuals** (e_t) are equal to the *difference between the observations and the corresponding fitted values*:

$$e_t = y_t - \hat{y}_t$$

Residual diagnostics

Residual diagnostics

Residuals are useful in checking whether a model has *adequately captured* the information in the data.

One thing that we **do not** want to observe in a forecasting model's residuals are **visible patterns**.

Thus, we need to learn some **tools** that can help us in exploring some possible patterns.

Residual diagnostics

Forecasting residuals should exhibit **two** main *properties*:

1. **No autocorrelation:**

- The presence of autocorrelation in a model's residuals implies that **useful information** that the model has ignored.

2. **Zero mean:**

- On average, what is left out of our explicit model should not have a value different than zero. Otherwise, our forecasting model may be **biased**.

In the **absence** of either (or both) of these properties, our task is to **improve** our forecasting model.

Residual diagnostics

In addition to the two aforementioned properties, two **additional** residual features are **desirable**:

1. **Constant variance (homoskedasticity);**
2. **Normally distributed residuals.**


Although not absolutely necessary, satisfying these two additional properties allows for better forecasting **prediction intervals**.

An example

An example

U.S. Dollars to Euro spot exchange rate

Jan 2023 – Feb 2024 (daily)



Source: U.S. Federal Reserve System.

An example

U.S. Dollars to Euro spot exchange rate, Jan 2023 – Feb 2024 (daily)

30-day forecast using the naïve method




Source: U.S. Federal Reserve System.

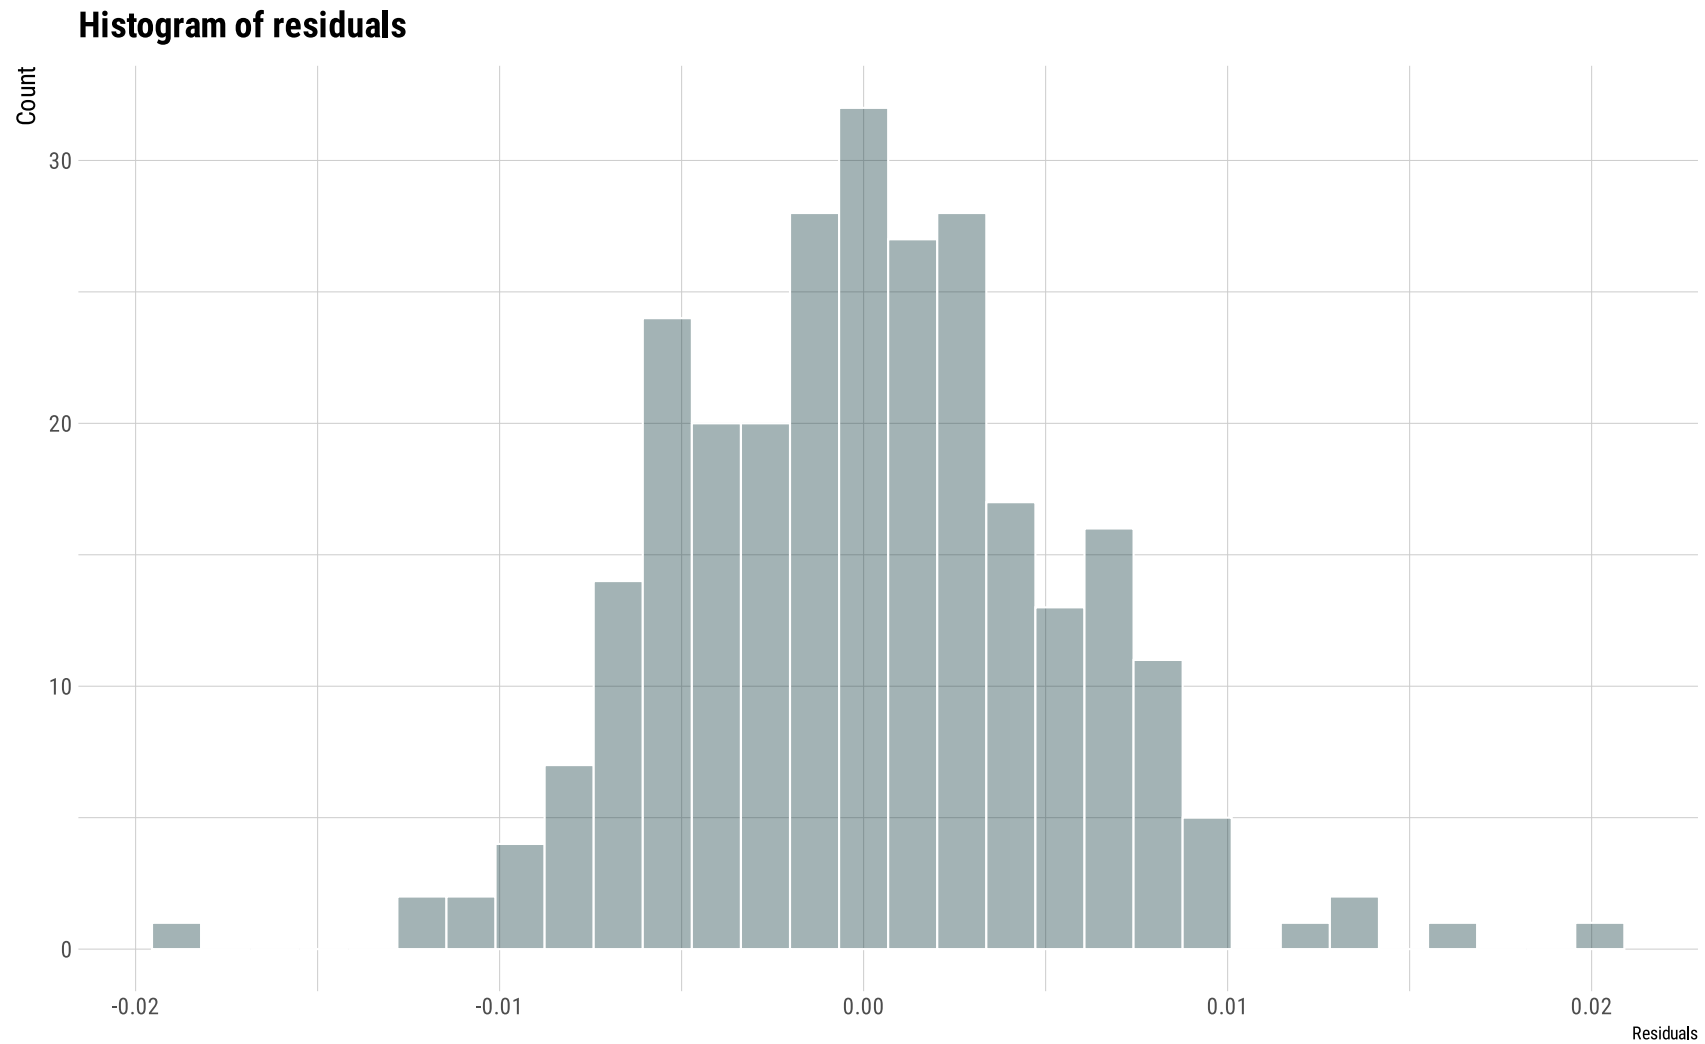
An example

```
#> # A tsibble: 290 x 6 [1]
#> # Key:      .model [1]
#>   .model      day  exch .fitted   .resid   .innov
#>   <chr>      <int> <dbl>   <dbl>   <dbl>   <dbl>
#> 1 naive_model 1016 NA      NA      NA      NA
#> 2 naive_model 1017 1.06   NA      NA      NA
#> 3 naive_model 1018 1.06   1.06   0.00610 0.00610
#> 4 naive_model 1019 1.05   1.06  -0.00980 -0.00980
#> 5 naive_model 1020 1.06   1.05   0.00970 0.00970
#> 6 naive_model 1021 1.08   1.06   0.0140  0.0140
#> 7 naive_model 1022 1.07   1.08  -0.00220 -0.00220
#> 8 naive_model 1023 1.07   1.07   0.000300 0.000300
#> 9 naive_model 1024 1.08   1.07   0.00830  0.00830
#> 10 naive_model 1025 1.08   1.08  -0.00120 -0.00120
#> # i 280 more rows
```


An example



An example



An example



Portmanteau tests

Portmanteau tests

Regarding residual autocorrelation, we can complement individual **ACF** analysis by using so-called **Portmanteau** tests.

The idea behind such tests is to consider a **whole set** of autocorrelation coefficients as a group, rather than testing for autocorrelation at individual lags, as we do in an ACF plot.

Thus, we can test whether the first l autocorrelations are significantly different from what would be expected from a **white noise** process.

When we test for autocorrelation in this grouped fashion, we call it a **Portmanteau** test.

We will further study **two** of these tests:

1. The *Box-Pierce* test;
2. The *Ljung-Box* test.

Portmanteau tests

- **The Box-Pierce test:**

The Box-Pierce test has the following **test statistic**:

$$Q = T \sum_{k=1}^l r_k^2$$

where T is the length of the time series; l is the number of lags you consider in the test; and r is the autocorrelation coefficient for each lag.

Regarding the **choice of lag** for this test, the textbook authors suggest $l = 10$ for *non-seasonal* data and $l = 2m$ for *seasonal* data (where m is the seasonal period).

Portmanteau tests

- **The Ljung-Box test:**

The Ljung-Box test has the following **test statistic**:

$$Q^* = T(T + 2) \sum_{k=1}^l (T - k)^{-1} r_k^2$$

Compared to the Box-Pierce test, it tends to give **more accurate** results.

For both tests, we assume a **Chi-squared** (χ^2) distribution with l degrees-of-freedom.

Portmanteau tests

As a **null hypothesis**, we assume that the *residuals are not distinguishable from a white noise series*.

For our previous example,

```
#> # A tibble: 1 × 3
#>   .model      bp_stat bp_pvalue
#>   <chr>      <dbl>    <dbl>
#> 1 naive_model  11.2      0.340
```

```
#> # A tibble: 1 × 3
#>   .model      lb_stat lb_pvalue
#>   <chr>      <dbl>    <dbl>
#> 1 naive_model  11.5      0.321
```

What do we conclude?

Prediction intervals

Prediction intervals

The inherent **uncertainty** from any forecast model is expressed by a **probability distribution**.

The usual way to display forecast results is by showing its (average) **point forecast**, along with a **prediction interval**.

A **prediction interval** gives an interval within which we expect y_t to lie with a specified probability.

For instance, a 95% prediction interval for an h -step forecast is given by

$$\hat{y}_{T+h|T} \pm 1.96\hat{\sigma}_h$$

where $\hat{\sigma}_h$ is the estimated **standard deviation** of the h -step forecast distribution.

Prediction intervals

In general, the prediction interval is obtained by:

$$\hat{y}_{T+h|T} \pm c\hat{\sigma}_h$$

where the value of c depends on the **confidence level** we assume for the prediction interval.

The usual values are 1.64 for **90%**; 1.96 for **95%**; and 2.58 for **99%**.

Prediction intervals

The estimated standard deviation parameter ($\hat{\sigma}_h$) will have *different* definitions depending on the forecasting **method** we adopt.

For our 4 **benchmark** methods:

Mean: $y_{T+h|T} \sim N(\bar{y}, (1 + 1/T)\hat{\sigma}^2)$

Naïve: $y_{T+h|T} \sim N(y_T, h\hat{\sigma}^2)$

Seasonal naïve: $y_{T+h|T} \sim N(y_{T+h-m(k+1)}, (k+1)\hat{\sigma}^2)$

Drift: $y_{T+h|T} \sim N(y_T + \frac{h}{T-1}(y_T - y_1), h\frac{T+h}{T}\hat{\sigma}^2)$

+

Prediction intervals

For our previous example,

```
#> # A tsibble: 30 x 6 [1]
#> # Key:      .model [1]
#>   .model      day      exch .mean      `80%`      `95%`
#>   <chr>      <dbl>    <dist> <dbl>    <hilo>      <hilo>
#> 1 naive_model 1306 N(1.1, 2.7e-05) 1.08 [1.071523, 1.084877]80 [1.067989, 1.088411]95
#> 2 naive_model 1307 N(1.1, 5.4e-05) 1.08 [1.068758, 1.087642]80 [1.063760, 1.092640]95
#> 3 naive_model 1308 N(1.1, 8.1e-05) 1.08 [1.066636, 1.089764]80 [1.060514, 1.095886]95
#> 4 naive_model 1309 N(1.1, 0.00011) 1.08 [1.064847, 1.091553]80 [1.057778, 1.098622]95
#> 5 naive_model 1310 N(1.1, 0.00014) 1.08 [1.063271, 1.093129]80 [1.055368, 1.101032]95
#> 6 naive_model 1311 N(1.1, 0.00016) 1.08 [1.061846, 1.094554]80 [1.053188, 1.103212]95
#> 7 naive_model 1312 N(1.1, 0.00019) 1.08 [1.060535, 1.095865]80 [1.051184, 1.105216]95
#> 8 naive_model 1313 N(1.1, 0.00022) 1.08 [1.059316, 1.097084]80 [1.049319, 1.107081]95
#> 9 naive_model 1314 N(1.1, 0.00024) 1.08 [1.058170, 1.098230]80 [1.047567, 1.108833]95
#> 10 naive_model 1315 N(1.1, 0.00027) 1.08 [1.057087, 1.099313]80 [1.045910, 1.110490]95
#> # i 20 more rows
```

Prediction intervals

U.S. Dollars to Euro spot exchange rate, Jan 2023 – Feb 2024 (daily)

30-day forecast using the naïve method



Source: U.S. Federal Reserve System.

Next time: Using transformations and decompositions