

ARIMA models: Modeling and forecasting

EC 361–001

Prof. Santetti

Spring 2024

Materials

Required readings:

- Hyndman & Athanasopoulos, ch. 9
 - sections 9.7–9.8.

Motivation

Motivation

Last time, we have joined the AR, I, and MA portions of **ARIMA** models.

When deciding on the **order** of our ARIMA models, a great starting point is to look at the *autocorrelation coefficient function* (**ACF**) and *partial autocorrelation function* (**PACF**) plots.

However, these can only do **so much** when deciding on which model to choose and perform our forecasts.

Therefore, we must move on to more **robust** procedures.

ARIMA at work

ARIMA at work


The `{fable}` R package handles ARIMA estimation following the **Hyndman-Khandakar algorithm**:

Hyndman-Khandakar algorithm for automatic ARIMA modelling

1. The number of differences $0 \leq d \leq 2$ is determined using repeated KPSS tests.
2. The values of p and q are then chosen by minimising the AICc after differencing the data d times. Rather than considering every possible combination of p and q , the algorithm uses a stepwise search to traverse the model space.
 - a. Four initial models are fitted:
 - $\text{ARIMA}(0, d, 0)$,
 - $\text{ARIMA}(2, d, 2)$,
 - $\text{ARIMA}(1, d, 0)$,
 - $\text{ARIMA}(0, d, 1)$.A constant is included unless $d = 2$. If $d \leq 1$, an additional model is also fitted:
 - $\text{ARIMA}(0, d, 0)$ without a constant.
 - b. The best model (with the smallest AICc value) fitted in step (a) is set to be the “current model”.
 - c. Variations on the current model are considered:
 - vary p and/or q from the current model by ± 1 ;
 - include/exclude c from the current model.The best model considered so far (either the current model or one of these variations) becomes the new current model.
 - d. Repeat Step 2(c) until no lower AICc can be found.


ARIMA at work

The **stepwise** procedure:




ARIMA at work

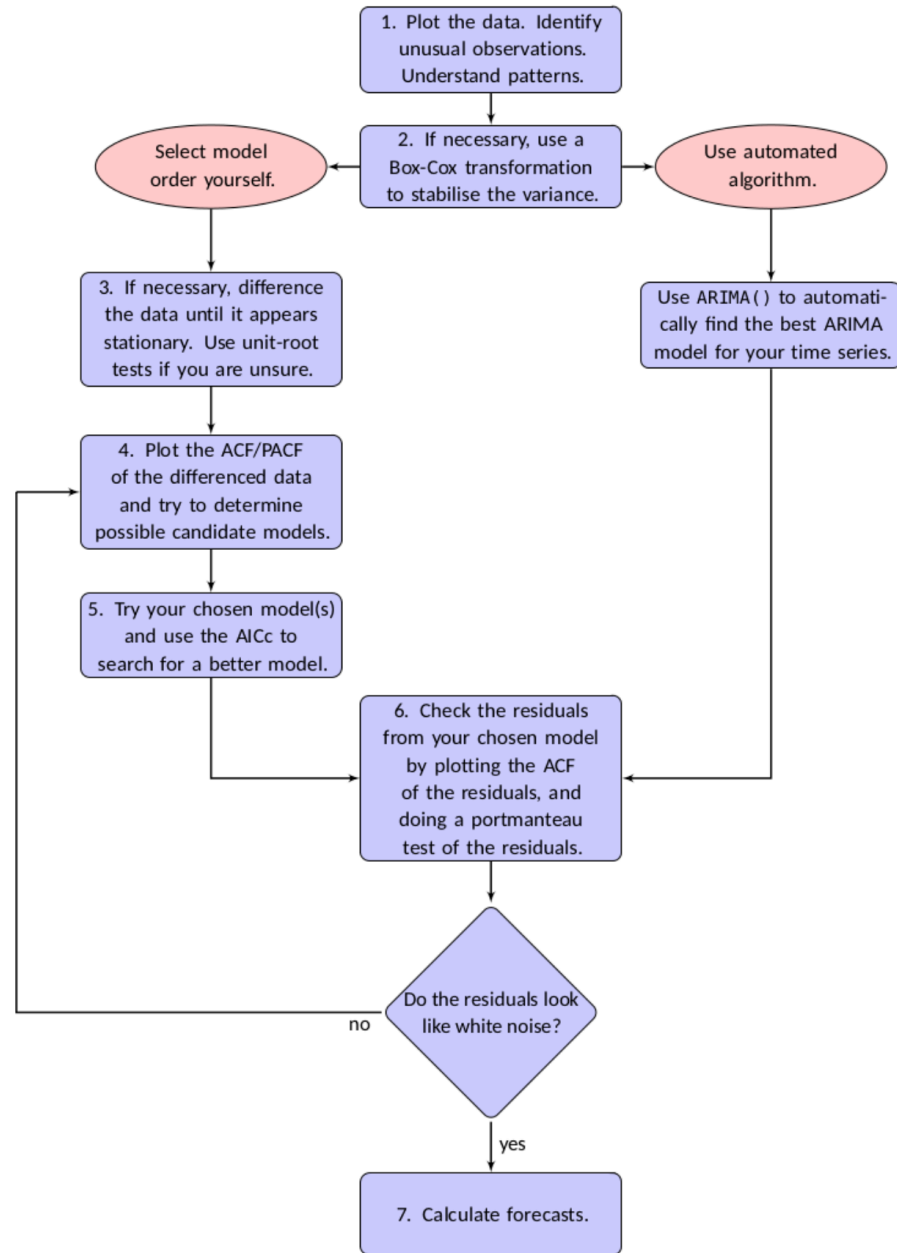
The **stepwise** procedure:



ARIMA at work

The **stepwise** procedure:






An example

An example

U.S. personal saving rate

1959–2023



Source: U.S. Bureau of Economic Analysis.

An example



```
saving_ts ▷  
  features(psav, unitroot_kpss)
```

```
#> # A tibble: 1 × 2  
#>   kpss_stat kpss_pvalue  
#>   <dbl>     <dbl>  
#> 1     1.22     0.01
```

```
saving_ts ▷  
  features(difference(psav), unitroot_kpss)
```

```
#> # A tibble: 1 × 2  
#>   kpss_stat kpss_pvalue  
#>   <dbl>     <dbl>  
#> 1   0.0537     0.1
```

An example



An example

```
saving_arima_fit <- saving_ts ▷  
  model(arima310 = ARIMA(psav ~ 1 + pdq(3, 1, 0)), # "1" includes a constant (c).  
        arima113 = ARIMA(psav ~ 1 + pdq(1, 1, 3)), # "1" includes a constant (c).  
        arima_auto = ARIMA(psav)) # letting {fable} select the best model.  
  
saving_arima_fit
```

```
#> # A mable: 1 x 3  
#>           arima310           arima113           arima_auto  
#>           <model>           <model>           <model>  
#> 1 <ARIMA(3,1,0) w/ drift> <ARIMA(1,1,3) w/ drift> <ARIMA(0,1,2)>
```




An example

```
saving_arima_fit >  
  glance() >  
  arrange(AICc) >  
  select(.model, AIC, AICc)
```

```
#> # A tibble: 3 × 3  
#>   .model      AIC  AICc  
#>   <chr>    <dbl> <dbl>  
#> 1 arima_auto 250.  250.  
#> 2 arima113   249.  250.  
#> 3 arima310   254.  255.
```



```
saving_arima_fit ▷  
  select(arima310) ▷  
  gg_tsresiduals()
```


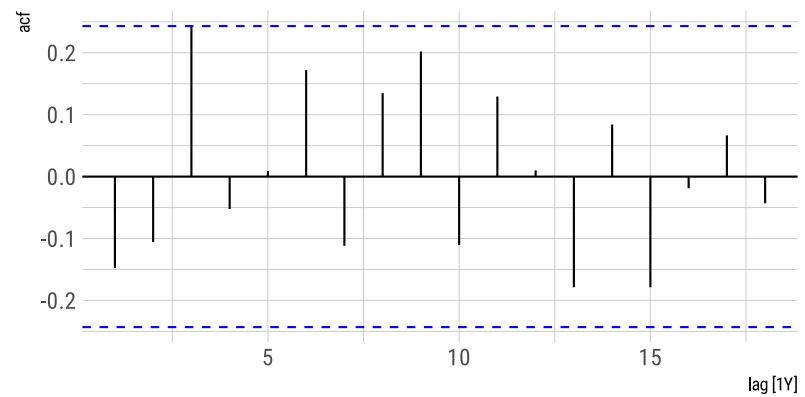



An example

```
saving_arima_fit ▷  
  augment() ▷  
  filter(.model = "arima310") ▷  
  features(.innov, ljung_box, lag = 10, dof = 3)
```

```
#> # A tibble: 1 × 3  
#>   .model    lb_stat lb_pvalue  
#>   <chr>      <dbl>    <dbl>  
#> 1 arima310    11.7      0.110
```

```
saving_arima_fit ▷  
  select(arima113) ▷  
  gg_tsresiduals()
```






An example

```
saving_arima_fit ▷  
  augment() ▷  
  filter(.model = "arima113") ▷  
  features(.innov, ljung_box, lag = 10, dof = 4)
```

```
#> # A tibble: 1 × 3  
#>   .model    lb_stat lb_pvalue  
#>   <chr>      <dbl>    <dbl>  
#> 1 arima113    15.2     0.0186
```

```
saving_arima_fit ▷  
  select(arima_auto) ▷  
  gg_tsresiduals()
```



An example

```
saving_arima_fit ▷  
  augment() ▷  
  filter(.model = "arima_auto") ▷  
  features(.innov, ljung_box, lag = 10, dof = 2)
```

```
#> # A tibble: 1 × 3  
#>   .model      lb_stat lb_pvalue  
#>   <chr>      <dbl>    <dbl>  
#> 1 arima_auto    8.42     0.393
```

Forecasting with ARIMA

Forecasting with ARIMA

Time to explain **how** ARIMA forecasts are generated.

Suppose our model of choice is an **ARMA(2, 2)**.

$$y_t = c + \phi_1 y_{t-1} + \phi_2 y_{t-2} + \theta_1 \varepsilon_{t-1} + \theta_2 \varepsilon_{t-2} + \varepsilon_t$$

Let us derive the **point forecasts** for $T + 1$ and $T + 2$:

Forecasting with ARIMA

Coming back to our example...

```
saving_arima_fc ← saving_arima_fit ▷  
  forecast(h = 6)
```

Forecasting with ARIMA

6-year ahead forecast: ARIMA(3, 1, 0) model

U.S. personal saving rate



Forecasting with ARIMA

6-year ahead forecast: ARIMA(1, 1, 3) model

U.S. personal saving rate



Forecasting with ARIMA

6-year ahead forecast: ARIMA(0, 1, 2) model

U.S. personal saving rate



Next time: Seasonal ARIMA models