

ARIMA models: Addressing seasonality

EC 361–001

Prof. Santetti
Spring 2024

Materials

Required readings:

- Hyndman & Athanasopoulos, ch. 9
 - sections 9.9—9.10.
 - reading 9.10 is *optional*.

Motivation

Motivation

By now, we know that time series that show prominent **trends** will *not* be stationary.

In addition, when **seasonality** is present, the time series **will not be stationary**.

There are specific **ARIMA** models that deal with seasonal series, and this is what this lecture is about.

Seasonal ARIMA models

Seasonal ARIMA models

A **seasonal ARIMA** model is formed by including **additional seasonal terms** in the ARIMA models we have seen so far.

A seasonal ARIMA model can be expressed as follows:

$$\text{ARIMA} \quad \underbrace{(p, d, q)}_{\text{Non-seasonal part}} \quad \underbrace{(P, D, Q)_m}_{\text{Seasonal part}}$$

Notice that we will adopt **upper-case** letters (P, D, Q) to express the ARIMA terms of the seasonal component.

Seasonal ARIMA models

For example, an **ARIMA(1, 1, 1) (2, 1, 1)₄** denotes a seasonal ARIMA model for *quarterly* data ($m = 4$), where to achieve stationarity we need:

- to take first differences ($d = 1$);
- to take an additional seasonal difference ($D = 1$).

In addition, the **seasonal** part has *two* autoregressive and *one* moving-average component, while the **non-seasonal** part has one AR and one MA component.

In terms of **order selection**, we still make use of **ACF** and **PACF** plots for help.

Seasonal ARIMA models

The seasonal part of an *ARIMA* model will be seen in the **seasonal lags** of the PACF and ACF plots.

As an example, consider an ***ARIMA*(0, 0, 0) (0, 0, 1)₁₂** model.

It will show the following features:

- A *spike* at lag 12 in the ACF but no other significant spikes;
- exponential *decay* in the seasonal lags of the PACF (i.e., at lags 12, 24, 36, ...).

In a similar way, for an ***ARIMA*(0, 0, 0) (1, 0, 0)₁₂** model, we will see:

- Exponential *decay* in the seasonal lags of the ACF;
- A single significant *spike* at lag 12 in the PACF.

Thus, when deciding on the *P* and *Q* values of a seasonal *ARIMA* model, we must restrict our attention to the **seasonal lags**.

Seasonal ARIMA models


The **Hyndman-Khandakar algorithm** looks and works basically in the same way as seen in the previous lecture.

It will only differ as additional parameters (namely P , D , and Q) need to be included in the mix.

Seasonal ARIMA models

International airline passengers


Jan 1949 – Dec 1960



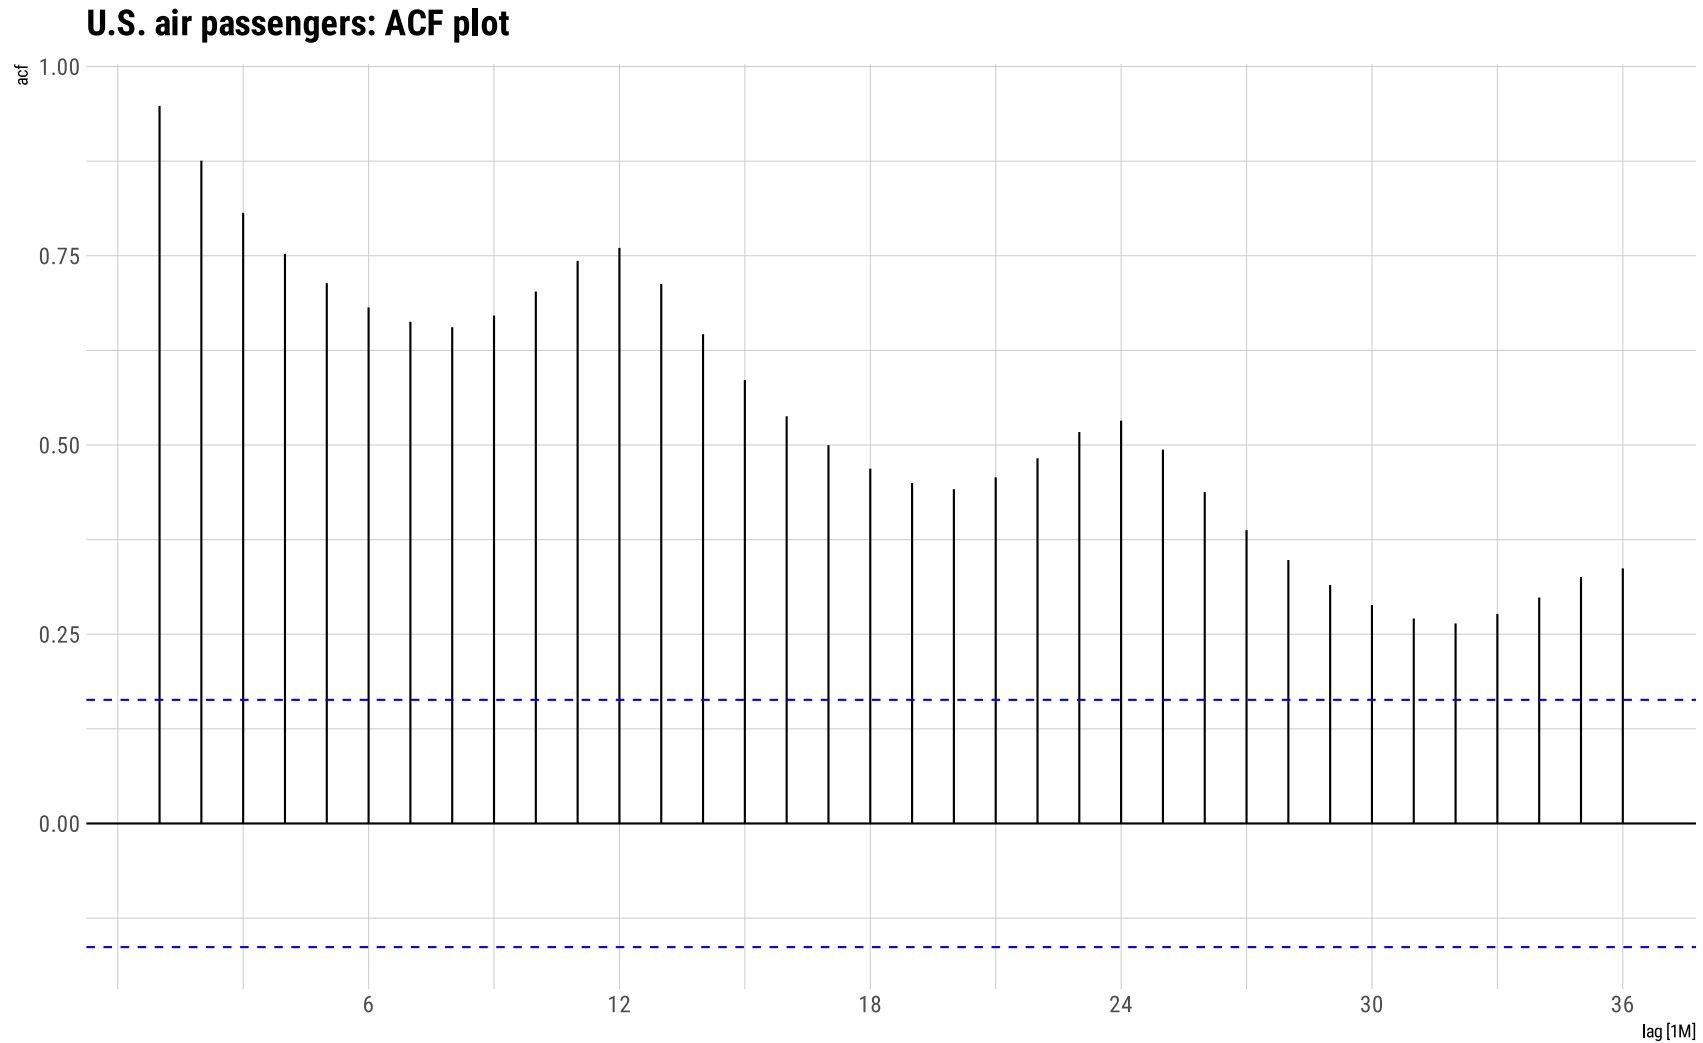
Source: Brown (1962).

Seasonal ARIMA models

U.S. air passengers: Seasonal plot



Seasonal ARIMA models



Seasonal ARIMA models

Is the series stationary?




```
air_ts >
  features(passengers, unitroot_kpss)

#> # A tibble: 1 × 2
#>   kpss_stat kpss_pvalue
#>   <dbl>     <dbl>
#> 1     2.74     0.01
```

```
air_ts ▷
```

```
mutate(seas_diff = difference(passengers, lag = 12)) ▷ # taking seasonal differencing.
```

```
gg_tsddisplay(seas_diff, plot_type = "partial")
```






Seasonal ARIMA models

Is the **seasonally differenced** series stationary?

```
air_ts >
  mutate(seas_diff = difference(passengers, lag = 12)) >
  features(seas_diff, unitroot_kpss)
```

```
#> # A tibble: 1 × 2
#>   kpss_stat kpss_pvalue
#>   <dbl>     <dbl>
#> 1    0.666    0.0167
```

```
air_ts >
  mutate(seas_diff = difference(passengers, lag = 12),
         dseas_diff = difference(seas_diff, lag = 1)) > #taking 1st difference of seasonal difference.
  gg_tsdisplay(dseas_diff, plot_type = "partial")
```



Seasonal ARIMA models

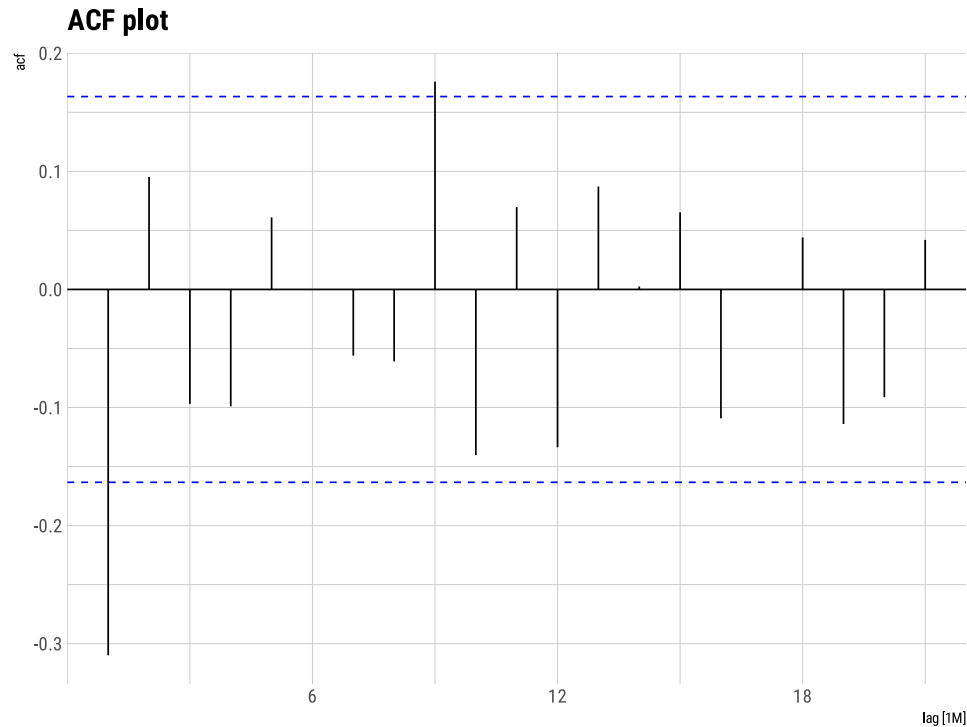
Have we achieved stationarity?

```
air_ts >
  mutate(seas_diff = difference(passengers, lag = 12),
         dseas_diff = difference(seas_diff, lag = 1)) >
  features(dseas_diff, unitroot_kpss)
```


```
#> # A tibble: 1 × 2
#>   kpss_stat kpss_pvalue
#>   <dbl>     <dbl>
#> 1  0.0543     0.1
```

Let us use the **ACF** and **PACF** plots to decide on the appropriate AR and MA orders.

```
air_ts >  
  ACF(dseas_diff) >  
  autoplot() +  
  labs(title = "ACF plot")
```



```
air_ts >  
  PACF(dseas_diff) >  
  autoplot() +  
  labs(title = "PACF plot")
```



Seasonal ARIMA models

Estimating a few models:

```
air_arima_fit <- air_ts ▷  
  model(arima110_010 = ARIMA(passengers ~ pdq(1, 1, 0) + PDQ(0, 1, 0)),  
        arima011_010 = ARIMA(passengers ~ pdq(0, 1, 1) + PDQ(0, 1, 0)),  
        arima310_010 = ARIMA(passengers ~ pdq(3, 1, 0) + PDQ(0, 1, 0)),  
        arima_auto = ARIMA(passengers))  
  
air_arima_fit
```



```
#> # A mable: 1 x 4  
#>           arima110_010           arima011_010           arima310_010           arima_auto  
#>           <model>           <model>           <model>           <model>  
#> 1 <ARIMA(1,1,0)(0,1,0)[12]> <ARIMA(0,1,1)(0,1,0)[12]> <ARIMA(3,1,0)(0,1,0)[12]> <ARIMA(2,1,1)(0,1,0)[12]>
```

Seasonal ARIMA models

```
air_arma_fit >
  glance() >
  arrange(AICc) >
  select(.model, AIC, AICc)
```

```
#> # A tibble: 4 × 3
#>   .model      AIC  AICc
#>   <chr>      <dbl> <dbl>
#> 1 arima_auto  1018. 1018.
#> 2 arima110_010 1020. 1020.
#> 3 arima011_010 1021. 1021.
#> 4 arima310_010 1024. 1024.
```

```
air_arma_fit >  
  select(arma_auto) >  
  gg_tsresiduals()
```



Seasonal ARIMA models

Are residuals **white noise**?


```
air_arma_fit >
  augment() >
  filter(.model = "arma_auto") >
  features(.innov, ljung_box, lag = 2 * 12, dof = 3)
```

```
#> # A tibble: 1 × 3
#>   .model      lb_stat lb_pvalue
#>   <chr>      <dbl>    <dbl>
#> 1 arma_auto  37.8      0.0137
```

Seasonal ARIMA models

International airline passengers (logs)


Jan 1949 – Dec 1960



Source: Brown (1962).

Seasonal ARIMA models

U.S. air passengers: Seasonal difference (logs)



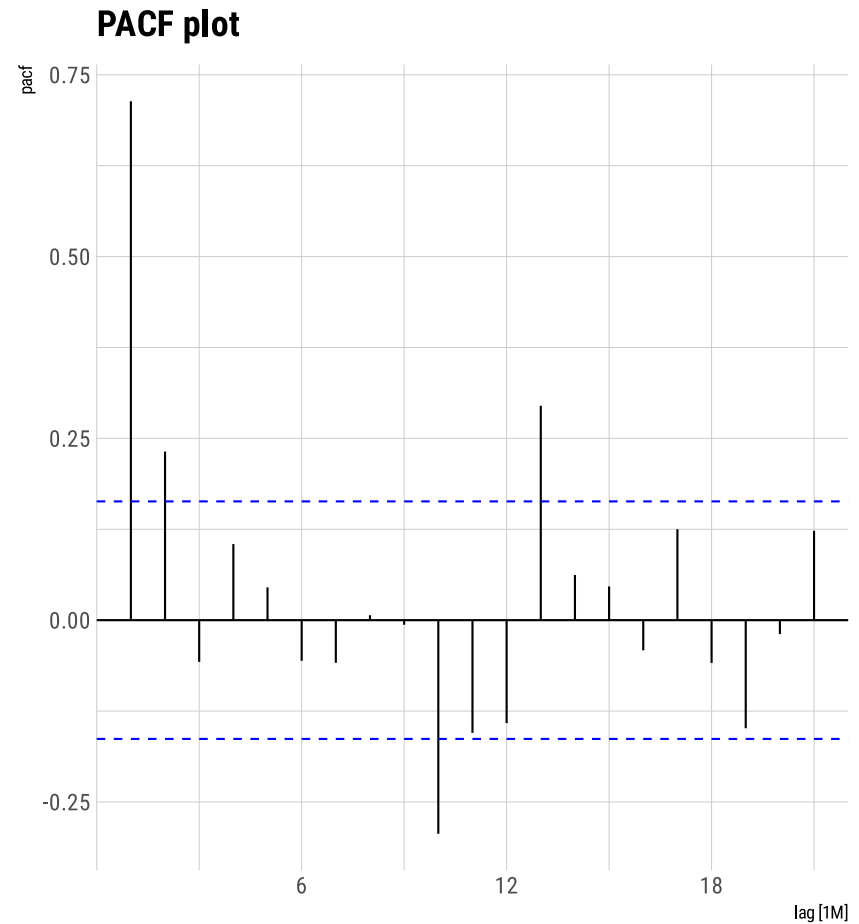

Seasonal ARIMA models

```
air_ts <- air_ts ▷  
  mutate(log_seas_diff = difference(log(passengers), lag = 12))
```

```
air_ts ▷  
  features(log_seas_diff, unitroot_kpss)
```

```
#> # A tibble: 1 × 2  
#>   kpss_stat kpss_pvalue  
#>   <dbl>     <dbl>  
#> 1    0.368    0.0909
```

Seasonal ARIMA models






Seasonal ARIMA models

```
air_log_fit <- air_ts ▷  
  model(arima200_011 = ARIMA(log(passengers) ~ 1 + pdq(2, 0, 0) + PDQ(0, 1, 1)),  
        arima202_011 = ARIMA(log(passengers) ~ 1 + pdq(2, 0, 2) + PDQ(0, 1, 1)),  
        arima_auto = ARIMA(log(passengers)))
```

```
air_log_fit ▷  
  glance() ▷  
  arrange(AICc) ▷  
  select(.model, AIC, AICc)
```

```
#> # A tibble: 3 × 3  
#>   .model      AIC  AICc  
#>   <chr>      <dbl> <dbl>  
#> 1 arima200_011 -489. -489.  
#> 2 arima_auto   -489. -489.  
#> 3 arima202_011 -486. -485.
```

```
air_log_fit >  
  select(arima_auto) >  
  gg_tsresiduals()
```



Seasonal ARIMA models

Are the residuals **white noise**?


```
air_log_fit >
  augment() >
  filter(.model = "arima_auto") >
  features(.innov, ljung_box, lag = 2 * 12, dof = 3)
```

```
#> # A tibble: 1 × 3
#>   .model      lb_stat lb_pvalue
#>   <chr>      <dbl>    <dbl>
#> 1 arima_auto  29.0      0.113
```

Seasonal ARIMA models

24-month ahead forecast

U.S. air passengers



Seasonal ARIMA models

```
air_log_fit >
  select(arima_auto) >
  forecast(h = 24) >
  head(5)
```

```
#> # A fable: 5 x 4 [1M]
#> # Key:      .model [1]
#>   .model      date      passengers .mean
#>   <chr>       <mth>      <dist>  <dbl>
#> 1 arima_auto 1961 Jan  t(N(6.1, 0.0013)) 453.
#> 2 arima_auto 1961 Feb  t(N(6.1, 0.0018)) 430.
#> 3 arima_auto 1961 Mar  t(N(6.2, 0.0022)) 486.
#> 4 arima_auto 1961 Apr  t(N(6.2, 0.0025)) 502.
#> 5 arima_auto 1961 May  t(N(6.3, 0.0028)) 522.
```

Next time: Dynamic regression models