

Dynamic regression models: Further thoughts

EC 361–001

Prof. Santetti

Spring 2024

Materials

Required readings:

- Hyndman & Athanasopoulos, ch. 10
 - sections 10.3; 10.5.

Motivation

Motivation

After studying the **essentials** of *dynamic regression* models, let us wrap up this topic with a few more **practical issues**.

Harmonic regression with ARIMA errors

Harmonic regression with ARIMA errors

We have already seen that **harmonic regression** (i.e., using *Fourier* terms) performs well for modeling **seasonality**.


- Especially *long* seasonal periods (*weekly* and *monthly* seasonal data, for example).

One **limitation**, however, is that harmonic regression assumes that the seasonal component is **fixed** over time (i.e., its variance is constant).

With this in mind, harmonic regression is a *useful tool* for seasonal time series.

Fourier terms can be used **along with** ARIMA errors to capture several important dynamics of the time series at hand.

Harmonic regression with ARIMA errors



Source: Brown (1962).

Harmonic regression with ARIMA errors

- Let us start a forecasting exercise for the **air passengers** data set with a *straightforward* regression model.




```
air_tslm <- air_ts ▷  
  model(reg = TSLM(log(passengers) ~ trend() + season()))
```



```
air_ts1m >
  report()
```

```
#> Series: passengers
#> Model: TSLM
#> Transformation: log(passengers)
#>
#> Residuals:
#>      Min       1Q   Median       3Q      Max
#> -0.156370 -0.041016  0.003677  0.044069  0.132324
#>
#> Coefficients:
#>              Estimate Std. Error t value Pr(>|t|)
#> (Intercept)    4.7267804   0.0188935 250.180 < 2e-16 ***
#> trend()         0.0100688   0.0001193  84.399 < 2e-16 ***
#> season()year2  -0.0220548   0.0242109  -0.911  0.36400
#> season()year3   0.1081723   0.0242118   4.468 1.69e-05 ***
#> season()year4   0.0769034   0.0242132   3.176  0.00186 **
#> season()year5   0.0745308   0.0242153   3.078  0.00254 **
#> season()year6   0.1966770   0.0242179   8.121 2.98e-13 ***
#> season()year7   0.3006193   0.0242212  12.411 < 2e-16 ***
#> season()year8   0.2913245   0.0242250  12.026 < 2e-16 ***
#> season()year9   0.1466899   0.0242294   6.054 1.39e-08 ***
#> season()year10  0.0085316   0.0242344   0.352  0.72537
#> season()year11 -0.1351861   0.0242400  -5.577 1.34e-07 ***
#> season()year12 -0.0213211   0.0242461  -0.879  0.38082
#> ---
#> Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
#>
#> Residual standard error: 0.0593 on 131 degrees of freedom
#> Multiple R-squared:  0.9835,    Adjusted R-squared:  0.982
#> F-statistic: 649.4 on 12 and 131 DF, p-value: < 2.22e-16
```

```
air_tsmlm >  
  augment() >  
  gg_tsdisplay(.innov, plot_type = "partial")
```



Harmonic regression with ARIMA errors

- Testing for **residual serial correlation**:

```
air_ts1m >  
  augment() >  
  features(.innov, ljung_box, lag = 2 * 12)
```

```
#> # A tibble: 1 × 3  
#>   .model lb_stat lb_pvalue  
#>   <chr>    <dbl>    <dbl>  
#> 1 reg      417.        0
```

What do we conclude?

Harmonic regression with ARIMA errors




Moving on to a model with **ARIMA errors**:

```
air_arima <- air_ts ▷  
  model(arima_errors = ARIMA(log(passengers) ~ trend() + season() + PDQ(0, 0, 0))) # setting PDQ() with zeros, as  
# the "season()" function is  
# taking care of seasonality.
```

```
air_arma >
  report()
```

```
#> Series: passengers
#> Model: LM w/ ARIMA(2,0,0) errors
#> Transformation: log(passengers)
#>
#> Coefficients:
#>      ar1      ar2  trend()  season()year2  season()year3  season()year4  season()year5  season()year6  season()year7  season()year8  season()year9  seas
#>      0.6746  0.1438    1e-02      -0.0211         0.1099         0.0794         0.0777         0.2004         0.3047         0.2958         0.1515
#> s.e.  0.0829  0.0836    4e-04         0.0106         0.0128         0.0145         0.0155         0.0161         0.0163         0.0162         0.0156
#>      season()year11  season()year12  intercept
#>             -0.1299             -0.0158         4.7282
#> s.e.             0.0131             0.0109         0.0306
#>
#> sigma^2 estimated as 0.001337:  log likelihood=279.55
#> AIC=-527.11  AICc=-522.82  BIC=-479.59
```

```
air_arima >  
  augment() >  
  gg_tsdisplay(.innov, plot_type = "partial")
```



Harmonic regression with ARIMA errors

- Testing for **residual serial correlation**:

```
air_arima ▷  
  augment() ▷  
  features(.innov, ljung_box, lag = 2 * 12)
```

```
#> # A tibble: 1 × 3  
#>   .model      lb_stat lb_pvalue  
#>   <chr>      <dbl>    <dbl>  
#> 1 arima_errors 36.6      0.0478
```

What do we conclude?

Harmonic regression with ARIMA errors


In case we forecast the time series at hand based on a model *with residual serial correlation*, **prediction intervals** will be **unreliable**.

```
air_arma >  
  forecast(h = 24) >  
  head(6)
```

```
#> # A tibble: 6 x 4 [1M]  
#> # Key:   .model [1]  
#>   .model      date      passengers .mean  
#>   <chr>      <mth>          <dbl> <dbl>  
#> 1 arima_errors 1961 Jan t(N(6.1, 0.0013)) 451.  
#> 2 arima_errors 1961 Feb t(N(6.1, 0.0019)) 450.  
#> 3 arima_errors 1961 Mar t(N(6.3, 0.0024)) 523.  
#> 4 arima_errors 1961 Apr t(N(6.2, 0.0028)) 516.  
#> 5 arima_errors 1961 May  t(N(6.3, 0.003)) 523.  
#> 6 arima_errors 1961 Jun t(N(6.4, 0.0032)) 600.
```



```
air_arma > forecast(h = 24) > autoplot(air_ts, level = 95, linewidth = .8) + labs(x = "")
```



Harmonic regression with ARIMA errors

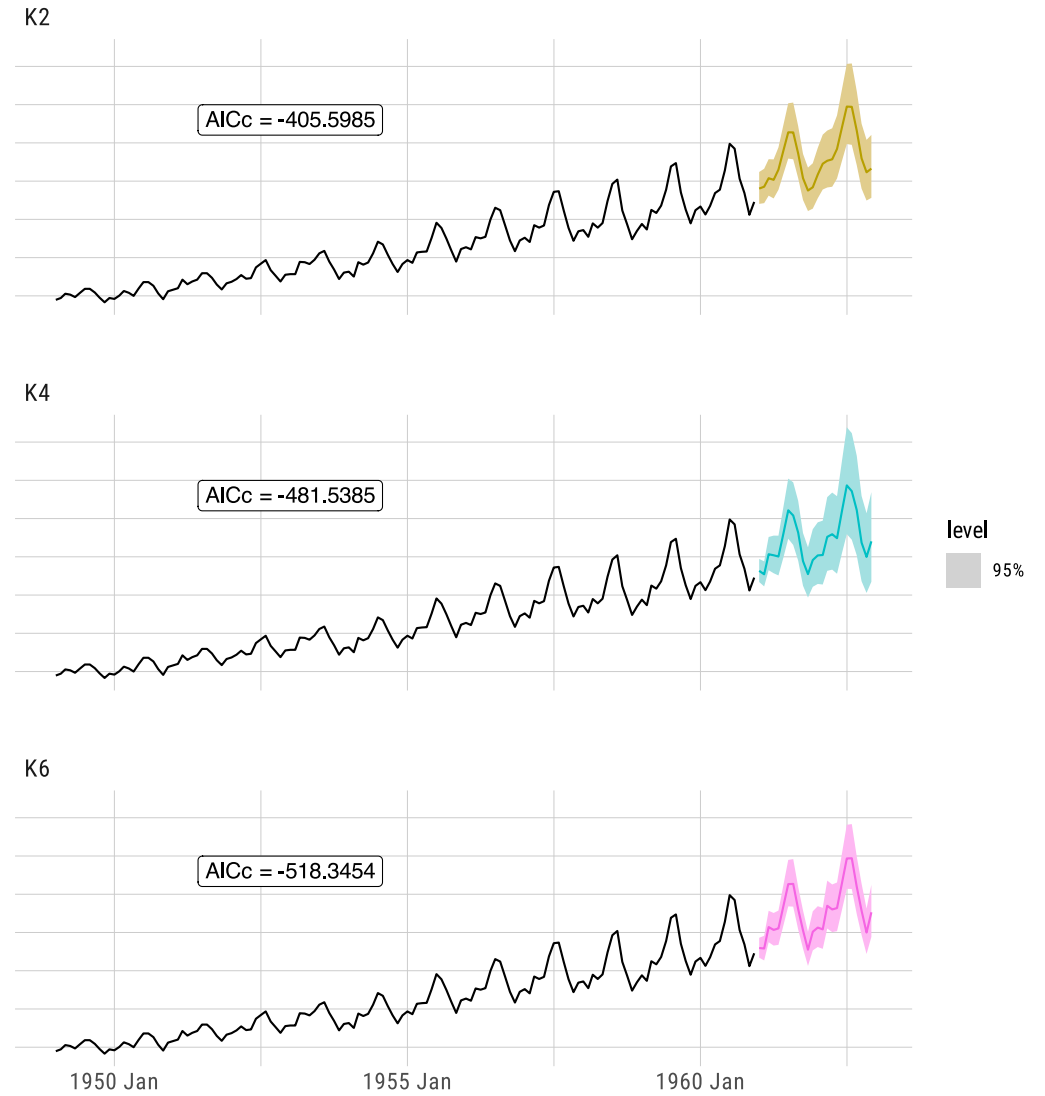

So far, we have not been successful in coming up with a *reliable* **regression forecast model** for the air passengers data set.

An *alternative* for modeling seasonality is using **harmonic regression with ARIMA errors**.



```
air_harmonic_fit <- model(air_ts,  
  K1 = ARIMA(log(passengers) ~ fourier(K = 1) + PDQ(0, 0, 0)),  
  K2 = ARIMA(log(passengers) ~ fourier(K = 2) + PDQ(0, 0, 0)),  
  K3 = ARIMA(log(passengers) ~ fourier(K = 3) + PDQ(0, 0, 0)),  
  K4 = ARIMA(log(passengers) ~ fourier(K = 4) + PDQ(0, 0, 0)),  
  K5 = ARIMA(log(passengers) ~ fourier(K = 5) + PDQ(0, 0, 0)),  
  K6 = ARIMA(log(passengers) ~ fourier(K = 6) + PDQ(0, 0, 0)))
```

Recall that K , the number of Fourier sine and cosine pairs, can vary from $K = 1$ up to $K = m/2$.

U.S. air passengers: 24-month ahead forecast



```
air_harmonic_fit >  
  select(K6) >  
  augment() >  
  gg_tsdisplay(.innov, plot_type = "partial")
```



Harmonic regression with ARIMA errors

```
air_harmonic_fit ▷  
  select(K6)
```

```
#> # A mable: 1 x 1  
#>                                K6  
#>                                <model>  
#> 1 <LM w/ ARIMA(1,1,1) errors>
```

```
air_harmonic_fit ▷  
  select(K5)
```

```
#> # A mable: 1 x 1  
#>                                K5  
#>                                <model>  
#> 1 <LM w/ ARIMA(1,1,1) errors>
```

```
air_harmonic_fit ▷  
  select(K6) ▷  
  augment() ▷  
  features(.innov, ljung_box, lag = 2 * 12, dof = 2)
```

```
#> # A tibble: 1 × 3  
#>   .model lb_stat lb_pvalue  
#>   <chr>    <dbl>    <dbl>  
#> 1 K6      35.6     0.0335
```


```
air_harmonic_fit ▷  
  select(K5) ▷  
  augment() ▷  
  features(.innov, ljung_box, lag = 2 * 12, dof = 2)
```

```
#> # A tibble: 1 × 3  
#>   .model lb_stat lb_pvalue  
#>   <chr>    <dbl>    <dbl>  
#> 1 K5      33.9     0.0503
```

Harmonic regression with ARIMA errors

24-month ahead forecast (Dynamic harmonic regression, $K = 5$)

U.S. air passengers data



Scenario-based forecasting

Scenario-based forecasting

When using **regression** models for forecasting, the time series we are interested in forecasting is a **function** of one or more **predictor variables**.

Excluding variables we **do know** their future values (e.g., *trend* and *seasonal dummy* variables), we must **make assumptions** about other predictors for which we **do not have** information.

This allows us to assume different **scenarios** for the *predictor variables* that are of interest.

Let us come back to our **Phillips curve** example.

Scenario-based forecasting

```
phillips_arima <- phillips_ts ▷  
  model(arima_reg = ARIMA(delta_infrate ~ unrata))
```

We will assume **3** possible **forecasting scenarios**:


1. The unemployment rate follows its historical *average value* (5.9%);
2. The unemployment rate *increases* to 10%;
3. The unemployment rate follows its *current value* of 3.36%.

Scenario-based forecasting


```
future_scenarios <- scenarios(average = new_data(phillips_ts, n = 5) ▷  
  mutate(unrate = mean(phillips_ts$unrate)),  
  pessimistic = new_data(phillips_ts, n = 5) ▷  
  mutate(unrate = 10),  
  optimistic = new_data(phillips_ts, n = 5) ▷  
  mutate(unrate = 3.36))
```

```
phillips_scen_fc <- phillips_arima ▷  
  forecast(h = 5, new_data = future_scenarios)
```


```
phillips_ts >
  autoplot(delta_infrate) +
  autolayer(phillips_scen_fc > filter(.scenario = "average"), level = 95, color = "#4682b4", linewidth = .8) +
  labs(title = "Average scenario")
```




```
phillips_ts >
  autoplot(delta_infrate) +
  autolayer(phillips_scen_fc > filter(.scenario = "pessimistic"), level = 95, color = "#800020", linewidth = .8) +
  labs(title = "Pessimistic scenario")
```



```
phillips_ts >
  autoplot(delta_infrate) +
  autolayer(phillips_scen_fc > filter(.scenario = "optimistic"), level = 95, color = "#228b22", linewidth = .8) +
  labs(title = "Optimistic scenario")
```



All 3 scenarios



Next time: Applications